

1 Model and Likelihood

Consider a linear K -factor model for D assets of the form

$$y_{it} = \alpha_d + \mathbf{f}_t' \boldsymbol{\beta}_d + \varepsilon_{it}$$

where $d = 1, \dots, D$ and $t = 1, \dots, T$ and $\mathbf{f}_t' = (f_{t1}, \dots, f_{tK})$ is a $K \times 1$ vector. This is a special case of the seemingly unrelated regression (SUR) model in which the regressors are *identical* across equations. Stacking observations for a given time period across assets, define $\mathbf{y}_t' = (y_{1t}, \dots, y_{Dt})$ and analogously $\boldsymbol{\varepsilon}_t' = (\varepsilon_{t1}, \dots, \varepsilon_{tD})$. Now let $\mathbf{x}_t' = (1, \mathbf{f}_t)$ and $\boldsymbol{\gamma}_d' = (\alpha_d, \boldsymbol{\beta}_d')$ so we have

$$\mathbf{y}_t = X_t \boldsymbol{\gamma} + \boldsymbol{\varepsilon}_t$$

where $X_t = I_D \otimes \mathbf{x}_t'$ and $\boldsymbol{\gamma}' = (\boldsymbol{\gamma}_1', \dots, \boldsymbol{\gamma}_D')$. Now, suppose that

$$\boldsymbol{\varepsilon}_t | \mathbf{x}_t \sim \text{iid } \mathcal{N}_D(0, \Sigma)$$

It follows that

$$L(\beta, \Sigma) \propto |\Sigma|^{-T/2} \exp \left[-\frac{1}{2} \sum_{t=1}^T (\mathbf{y}_t - X_t \boldsymbol{\gamma})' \Sigma^{-1} (\mathbf{y}_t - X_t \boldsymbol{\gamma}) \right]$$

We parameterize this problem in terms of the $D \times D$ precision matrix Σ and the $D(K+1) \times 1$ vector of regression coefficients $\boldsymbol{\gamma}$.

2 Prior and Posterior Distribution

To complete the model we specify a prior distribution on the full collection of parameters $\boldsymbol{\theta} = (\boldsymbol{\gamma}, \Omega^{-1})$, namely

$$\pi(\boldsymbol{\theta}) = \mathcal{N}_{d+p}(\boldsymbol{\gamma} | \boldsymbol{\gamma}_0, G_0) \mathcal{W}_d(\Omega^{-1} | \rho_0, R_0)$$

This prior is conditionally conjugate with the normal likelihood yielding the posterior distribution $Y_n = (\{y_t\}, \{f_t\})$

$$\pi(\boldsymbol{\theta} | Y_n) \propto \mathcal{N}_{d+p}(\boldsymbol{\gamma} | \boldsymbol{\gamma}_0, G_0) \mathcal{W}_d(\Omega^{-1} | \rho_0, R_0) \prod_{t=1}^n \mathcal{N}_d(y_t | X_t \boldsymbol{\gamma}, \Omega)$$

3 MCMC

We can summarize this posterior distribution by the MCMC algorithm of Chib and Greenberg (1995). It requires the repeated recursive sampling of the following two conditional distributions:

$$\gamma|Y_n, \Omega^{-1} \sim \mathcal{N}_{d+p}(\hat{\gamma}, G_n)$$

where

$$\begin{aligned}\hat{\gamma} &= G_n \left(G_0^{-1} \gamma_0 + \sum_{t=1}^n X_t' \Omega^{-1} y_t \right) \\ G_n &= \left(G_0^{-1} + \sum_{t=1}^n X_t' \Omega^{-1} X_t \right)^{-1}\end{aligned}$$

and

$$\Omega^{-1}|Y_n, \gamma \sim \mathcal{W}_d \left(\rho_0 + n, \left(R_0^{-1} + \sum_{t=1}^n (y_t - X_t \gamma) (y_t - X_t \gamma)' \right)^{-1} \right)$$

4 Marginal likelihood

The marginal likelihood is available by the method of Chib (1995). From the Chib (1995) identity, we have

$$\log m(\{y_t\}) = \log \pi(\gamma^*) + \log \pi(\Omega^{-1*}) + \sum_{t=1}^n \log p(y_t | X_t \gamma^*, \Omega^*) - \log \pi(\gamma^*, \Omega^{-1*} | Y_n)$$

where the last term is calculated as

$$\pi(\Omega^{-1*} | Y_n) \times \pi(\gamma^* | Y_n, \Omega^{-1*})$$

in which the first term is estimated by averaging the full-conditional Wishart density over the draws $\{\gamma^{(g)}\}_{g=1}^G$ from the main MCMC run

$$\pi(\Omega^{-1*} | Y_n) = \frac{1}{G} \sum_{g=1}^G \mathcal{W}_d \left(\Omega^{-1*} | \rho_0 + n, \left(R_0^{-1} + \sum_{t=1}^n (y_t - X_t \gamma^{(g)}) (y_t - X_t \gamma^{(g)})' \right)^{-1} \right)$$

and the second term $\pi(\gamma^*|Y_n, \Omega^{-1*})$ is available directly as

$$\pi(\gamma^*|Y_n, \Omega^{-1*}) = \mathcal{N}_{d+p}(\gamma^*|\hat{\gamma}^*, G_n^*)$$

where

$$\begin{aligned}\hat{\gamma}^* &= G_n^* \left(G_0^{-1} \gamma_0 + \sum_{t=1}^n X_t' \Omega^{-1*} y_t \right) \\ G_n^* &= \left(G_0^{-1} + \sum_{t=1}^n X_t' \Omega^{-1*} X_t \right)^{-1}\end{aligned}$$

5 Prediction

Suppose we are interested in predicting the cross-section of returns y_{n+1} at time $(n+1)$. The Bayes prediction density of these returns, conditioned on the data Y_{n+1} and the factors f_{n+1} , is given by

$$p(y_{n+1}|Y_n, f_{n+1}) = \int_{\gamma, \Omega^{-1}} \mathcal{N}_d(y_{n+1}|X_{n+1}\gamma, \Omega) d\pi(\gamma, \Omega^{-1}|Y_n)$$

which is estimated by the ergodic Monte Carlo average

$$p(y_{n+1}|Y_n, f_{n+1}) = \frac{1}{G} \sum_{g=1}^G \mathcal{N}_d(y_{n+1}|X_{n+1}\gamma^{(g)}, \Omega^{(g)})$$

with the MCMC draws $\{\gamma^{(g)}, \Omega^{(g)}\}$ from the posterior distribution.

6 Student-t errors

Suppose now that the errors are distributed as multivariate-t

$$\varepsilon_t \sim t_{d,\nu}(0, \Omega)$$

so that

$$E(\varepsilon_t) = 0, \quad \nu > 1$$

$$Var(\varepsilon_t) = \frac{\nu}{\nu - 2} \Omega, \quad \nu > 2$$

The analysis of this model utilizes the hierarchical representation

$$\begin{aligned} \varepsilon_t | \lambda_t &\sim N(0, \lambda_t^{-1} \Omega) \\ \lambda_t &\sim G\left(\frac{\nu}{2}, \frac{\nu}{2}\right) \end{aligned}$$

which means that conditioned on $(\nu, \{\lambda_t\})$, the results presented for the Gaussian model can be applied with minor modifications. The MCMC sampling is completed with the sampling of $(\nu, \{\lambda_t\})$.

Following Albert and Chib (1993), let us assume that the support of ν is the set of values $\{\nu_j\}_{j=1}^J$, for example, $\{4, 6, 8, 10, 12, 14, 16\}$ and that a priori

$$\Pr(\nu = \nu_j) = q_j$$

Then, simple calculations show that

$$\gamma | Y_n, \Omega^{-1}, \nu, \{\lambda_t\} \sim \mathcal{N}_{d+p}(\hat{\gamma}_\lambda, G_{n,\lambda})$$

where

$$\begin{aligned} \hat{\gamma}_\lambda &= G_{n,\lambda} \left(G_0^{-1} \gamma_0 + \sum_{t=1}^n \lambda_t X_t' \Omega^{-1} y_t \right) \\ G_{n,\lambda} &= \left(G_0^{-1} + \sum_{t=1}^n \lambda_t X_t' \Omega^{-1} X_t \right)^{-1} \end{aligned}$$

and

$$\Omega^{-1} | Y_n, \gamma, \nu, \{\lambda_t\} \sim \mathcal{W}_d \left(\rho_0 + n, \left(R_0^{-1} + \sum_{t=1}^n \lambda_t (y_t - X_t \gamma) (y_t - X_t \gamma)' \right)^{-1} \right)$$

Moreover,

$$\Pr(\nu = \nu_j | Y_n, \gamma, \Omega^{-1}) \propto q_j \prod_{t=1}^n t_{d,\nu_j}(y_t | X_t \gamma, \Omega)$$

and

$$\lambda_t | Y_n, \gamma, \nu \sim G\left(\frac{\nu + d}{2}, \frac{\nu + (y_t - X_t \gamma)' (y_t - X_t \gamma)}{2}\right)$$

One sweep of the MCMC sampling is completed by sampling these four distributions in this order.

6.1 Marginal likelihood

The Chib (1995) method can again be applied to find the log marginal likelihood as

$$\log \Pr(\nu^*) + \log \pi(\gamma^*) + \log \pi(\Omega^{-1*}) + \sum_{t=1}^n \log t_{d,\nu^*}(y_t | X_t \gamma^*, \Omega^*) - \log \pi(\nu^*, \gamma^*, \Omega^{-1*} | Y_n)$$

where ν^* is the posterior mode (which is easily computed from the sampled values), the last term is calculated as

$$\Pr(\nu^* | Y_n) \times \pi(\Omega^{-1*} | Y_n, \nu^*) \times \pi(\gamma^* | Y_n, \Omega^{-1*}, \nu^*)$$

in which the first term is obtained from the posterior frequency distribution of ν , the second term is obtained from a reduced run in which ν is fixed at ν^* and the remaining three distributions are sampled and the draws

$$\left\{ \gamma^{(g)}, \lambda_t^{(g)} \right\}_{g=1}^G$$

from this reduced MCMC run are used to calculate $\pi(\Omega^{-1*} | Y_n, \nu^*)$ as

$$\frac{1}{G} \sum_{g=1}^G \mathcal{W}_d \left(\Omega^{-1*} | \rho_0 + n, \left(R_0^{-1} + \sum_{t=1}^n \lambda_t^{(g)} (y_t - X_t \gamma^{(g)}) (y_t - X_t \gamma^{(g)})' \right)^{-1} \right)$$

and the final term $\pi(\gamma^* | Y_n, \Omega^{-1*})$ is obtained from a second reduced run in which ν is fixed at ν^* and Ω^{-1} is fixed at Ω^{-1*} and the draws

$$\left\{ \lambda_t^{(g)} \right\}$$

from this reduced run are used to give

$$\pi(\gamma^* | Y_n, \Omega^{-1*}, \nu^*) = \frac{1}{G} \sum_{g=1}^G \mathcal{N}_{d+p} \left(\gamma^* | \hat{\gamma}_{\lambda^{(g)}}^*, G_{n,\lambda^{(g)}}^* \right)$$

where

$$\begin{aligned} \hat{\gamma}_{\lambda^{(g)}}^* &= G_{n,\lambda^{(g)}}^* \left(G_0^{-1} \gamma_0 + \sum_{t=1}^n \lambda_t^{(g)} X_t' \Omega^{-1*} y_t \right) \\ G_{n,\lambda^{(g)}}^* &= \left(G_0^{-1} + \sum_{t=1}^n \lambda_t^{(g)} X_t' \Omega^{-1*} X_t \right)^{-1} \end{aligned}$$